# Instrument, Object, Agent: Large Language Models and the Description of Variation and Change in World Englishes

**Jasper Kyle Catapang**[1]

[1] *Graduate School of Global Studies, Tokyo University of Foreign Studies, 3-11-1 Asahi-cho, Fuchu, Tokyo, Japan*

This paper argues that large language models (LLMs) stand in three relations to World Englishes research: as instruments that describe, generate and annotate data; as objects whose output is a describable, standard-leaning written register; and as agents that take part in communicative ecologies and supply norms. Extending earlier work on default model behavior, it tests four models of one family against benchmarks from Peter Collins's corpus studies of modality and the progressive. Describing Philippine, Singapore, Indian and Kenyan English, the models restated the literature's summaries, credited claims to Collins mostly correctly but with errors that leaned toward familiar stereotypes, and often tied local grammar to low proficiency. Writing local news, they localized vocabulary but not grammar, produced no stative progressives, all but removed *must* from British news and omitted Philippine *shall* and *might*. Annotating modal meaning, they were highly self-consistent, with an accuracy that depended on whose definitions they were given. The paper proposes the notion of algorithmic exonormativity and closes with a reporting checklist.

**Date:** November 14, 2026
**Correspondence:** catapang.jasper.kyle.y0@tufs.ac.jp
**ORCID:** https://orcid.org/0000-0002-4510-0975
**Website:** https://www.tufs.ac.jp/english/research/

東京外国語大学
Tokyo University of Foreign Studies

**Keywords:** large language models; World Englishes; modality; progressive aspect; algorithmic exonormativity; corpus linguistics; annotation; Peter Collins

## 1 Introduction

For more than three decades, the description of English worldwide has rested on a simple and demanding premise: claims about a variety must answer to what its speakers and writers actually produce. Few scholars have honored that premise as consistently as Peter Collins, whose corpus studies of the modals, the quasi-modals and the progressive across Inner and Outer Circle Englishes (Collins, 2008, 2009, 2023; Collins et al., 2014; on the circles, Kachru, 1985) turned loose talk of "conservative" and "innovative" varieties into measurable propositions. It is fitting, if a little uncomfortable, to ask in his honor what becomes of that premise now that fluent English can be produced on demand by large language models (LLMs) trained on much of the written record, including much of what has been written about Collins's findings. The allure is plain; the question is what, if anything, their output tells us about English as it is used in real communities.

LLMs stand in three relations to that enterprise. They are *instruments*, offered as heuristics, annotators and generators of data; *objects*, whose output has regularities of its own (Grieve et al., 2025; Reinhart et al., 2025); and *agents*, which take part in the communicative lives of speakers, from migrants handling bureaucracies (Catapang et al., 2025) to students drafting essays, and so act as normative infrastructure (Catapang, 2026). The relations differ in what a model's text is taken to be: a *measurement* of language, a *specimen* of it or a *cause* of it. Figure 1 maps the three relations and the error that arises when any two are conflated; the news articles generated for this study occupy all three at once.

The paper is a sequel to Catapang (2026), which found that three commercial systems treated Asian Englishes by default with *asymmetrical pluralism*, many Englishes present but unevenly authorized, and left open whether prompting can elicit variety features. Against Collins's baselines, the present

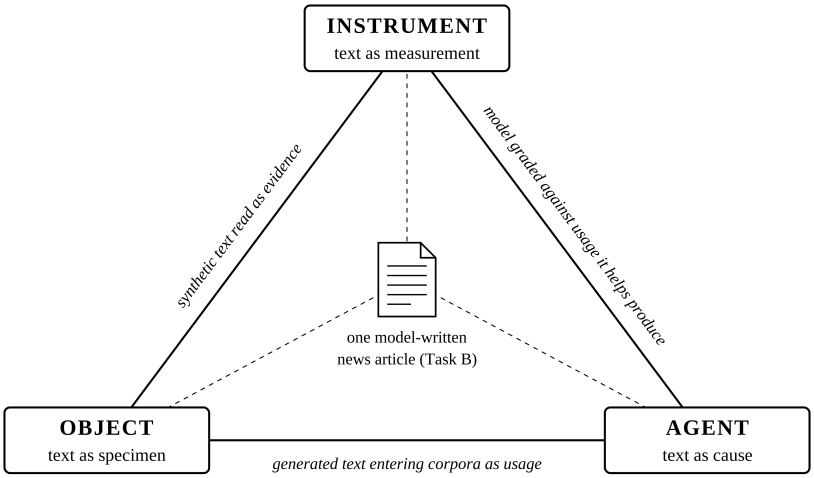



Figure 1: Three relations in which a language model's text can stand to World Englishes research. Each side names the error that results when the relations at its ends are conflated; a single news article generated for this study (Task B, Section 3) occupies all three.

paper asks whether the models *know* what the corpora show, *do* it when writing for a local readership, and *annotate* modal meaning as a trained analyst would. In describing, they reproduce the literature rather than the language, or more exactly the literature's summaries rather than its tables; in writing, they reproduce it only in part. After the evidentiary standard and design (Sections 2 and 3) and the results (Section 4), Section 5 reads the findings through the three relations, developing the notion of *algorithmic exonormativity*, and Section 6 offers a reporting checklist.

## 2 Instruments and evidence

Collins's program is an exacting benchmark because it is at once comparative, diachronic and explicit about numbers. His studies of the progressive (Collins, 2008) and of the modals and quasi-modals (Collins, 2009) draw on matched samples of nine varieties, among them British, American, Philippine, Singapore, Indian and Kenyan English, so that differences can more safely be attributed to variety than to sampling. Comparing written Philippine English of around 1960 and the early 1990s, Collins et al. (2014) found six lower-frequency modals declining more slowly ($-15.5\%$) than in British ($-21.1\%$) or American English ($-25.1\%$), concluded that Philippine English follows neither closely, though it was catching up with American English in some areas, and read this as endonormative stabilization, a placement that Borlongan (2016), a co-author, developed across Schneider's parameters. In the Global Web-based English Corpus (GloWbE; Davies & Fuchs, 2015), Outer Circle varieties use more *must* per million words (Philippines 633, Kenya 620, India 571, Singapore 527) than American (488) and British English (487), but not uniformly less *need to*: India (802), Kenya (708) and Singapore (691) exceed the United States (643), and only the Philippines (594) falls clearly below it (Collins, 2023). Collins's conclusion, that the Inner Circle is further advanced in the shift from modals to quasi-modals, holds for the system but not for every item.

Collins (2008) found the progressive twice as frequent in speech as in writing and ranked Australian and New Zealand English as the most innovative in its use, the Southeast Asian varieties next, American and British English after them, and Indian and Kenyan English last, while cautioning that the results were not conclusive. On stative progressives, among the most often cited features of the New Englishes, the evidence is split: van Rooy (2014) found them proportionally more frequent in Indian, Kenyan and Black South African than in British English, and Collins (2015) found them increasing in written Philippine English, whereas Rautionaho (2022) could not confirm higher Outer Circle rates. A feature so salient and so insecurely quantified tests whether models track salience rather than frequency.

On LLMs as *instruments*, enthusiasm for cheap annotation (Gilardi et al., 2023) has met caution about validation and prompt stability (Törnberg, 2024): a supervised pipeline exceeded 90% accuracy

on a binary grammatical variable (Morin & Marttinen Larsson, 2025), but a model's "frequency lists" overlapped only partly with corpus-derived ones (Uchida, 2024); in the volume for which this paper was written (hereafter the Collins volume), Ike and Fujiwara (Chapter 16) use an LLM to extract collocations. As *objects*, LLM outputs differ from human writing in their lexico-grammatical profiles (Reinhart et al., 2025). As *agents*, models stereotype and condescend to speakers of Indian, Nigerian, Singaporean, Kenyan and other Englishes (Fleisig et al., 2024), encode covert prejudice against speakers of African American English that alignment conceals (Hofmann et al., 2024), and through preference tuning widen the gap between American English and Indian and Nigerian English (Ryan et al., 2024). Catapang (2026), analyzing 240 interactions with GPT-5.1, Claude Opus 4.5 and Gemini 3 Pro, found a Singapore English *Can or not?* corrected to *Is it possible?* and metalinguistic answers that recognized Outer Circle Englishes while confining them to informal use: the *medium audibility* produced by what that study calls *differential audibility*. The literature has shown that models talk about Englishes and talk down to their speakers; it has not asked, against Collins's counts, whether they can write in them.

## 3 Design

Four models of one family were tested: claude-haiku-5-5, claude-sonnet-5-5 and claude-opus-5-5, three tiers of the current generation, and claude-fable-5-1, of the preceding one (hereafter Haiku, Sonnet, Opus and Fable), all on 8 October 2026. Each trial was a fresh session, run as a subagent in an agentic harness (Claude Agent SDK) and told to answer from its own knowledge; besides the prompt, it received the harness's system prompt, environment information and a file-writing tool, and sampling could not be set. The session context of every trial also named the project, describing that volume as a festschrift for Peter Collins, and relayed a message from the author's session naming his two papers on artificial intelligence and Englishes. No trace of either appears in the outputs, but the context bears on Task A (Section 4.1). Models from one vendor share a training pipeline, which licenses no generalization to GPT or Gemini systems.

The four Outer Circle varieties appear in Collins's matched corpora, and his findings separate them: the Southeast Asian varieties lead the Outer Circle in the shift to quasi-modals (Collins, 2009), and Indian and Kenyan English were the least innovative in their use of the progressive (Collins, 2008). A variety-blind model would erase these internal contrasts; a literature-driven one would reproduce whichever contrasts the literature makes most audible. News holds topic constant, has a human benchmark in the Philippine press (Dita & Chan Bona, Chapter 16, Collins volume), and produced the most distinctive result of Collins's diachronic work on Philippine English, press modals rising while British and American ones fell (Collins et al., 2014). Table 1 summarizes the tasks.

The Task A prompt named neither Collins nor any corpus, but it selected the four forms Collins (2009, 2023) analyzes, framed the task as a comparison with American and British English, and asked "whether and how" each variety extends the progressive, which presupposes the feature. The coders were instances of a model under test, and their dossier on Collins drew on publisher pages, abstracts and citing studies, since the per-variety figures of Collins (2008, 2009) were not accessible. Same-family agreement shows consistency rather than validity (Section 4.3), hence the second, independently coded pass.

The Task B prompt named a country and a local readership but not a variety, so it tests situated rather than feature-targeted elicitation. The 20 topics were identical across countries but follow a Philippine news agenda (an approaching typhoon, dengue, record remittances), so topic and variety are partly confounded. Because tokens cluster within articles and topics, only Outer–Inner Circle contrasts that survive both Holm correction and the two clustering checks are treated as findings. GloWbE serves for direction and ratio; the 2022 Philippine English Press Corpus (PEPC) supplies register-matched levels (*must* 32.8, *have to* 50.2, *need to* 30.3, *should* 66.2 per 100,000 words), although its operational definitions may differ from those used here.

Task C used 135 examples, covering 32 modal and quasi-modal forms, from a study of modality in the Philippine press (Chan Bona, 2025) that classifies meanings within the framework of Huddleston and Pullum (2002). Each gold label is the one implied by the section of the thesis that presents the example (36 deontic, 59 epistemic, 40 dynamic), as extracted independently by two model instances with complete agreement on 152 candidates, which checks transcription, not validity.

The generated articles are fiction, analyzed only as model output and never pooled with attested data;

Table 1: Design of the three tasks

| Task | Prompt | Trials and output | Coding and analysis | Benchmark |
|---|---|---|---|---|
| A: Know | Describe how the variety uses *must*, *have to*, *need to* and *should*, with any relative frequencies, and "whether and how" it uses the progressive with stative verbs, compared with American and British English | 4 models × 4 varieties (PH, SG, IN, KE) × 2 repetitions = 32 descriptions | Two independent claude-opus-5-5 instances, given the rubric and anonymized texts but not the hypotheses (agreement 90.6–100%, $\kappa$ = .762–1.000); 44 sentences naming Collins coded against a dossier of his published findings (72.7%, $\kappa$ = .574); framing and register fields double-coded in a second pass (agreement 96.9–100%, $\kappa$ = .949–1.000) | Collins (2008, 2009, 2015, 2023); Collins et al. (2014); van Rooy (2014); Rautionaho (2022) |
| B: Do | Ten news articles of about 450 words each, as published in a national daily in the named country, by local journalists for a local readership, on 20 fixed topics; no variety named | 4 models × 6 countries (PH, SG, IN, KE, US, GB) × 2 batches of 10; 450 articles after refusals, 224,186 words | Automatic counts; stative progressives sought with 18 and with about 45 verbs (optional adverb), hits read; rates per 100,000 words; rate ratios with exact 95% CIs; log-likelihood (LL), Holm-corrected across four forms; 95% CIs from resampling the 20 topic slots; permutation of variety labels within model-by-topic strata | GloWbE (Collins, 2023); PEPC 2022 (Dita & Chan Bona, Chapter 16, Collins volume) |
| C: Annotate | Classify a marked modal expression as deontic, epistemic, dynamic or other, following short definitions with examples | 4 models × 2 runs × 135 items = 1,080 labels | Accuracy with 95% CIs; Cohen's $\kappa$ with gold and between runs; exact McNemar tests, Holm-corrected across 12 comparisons | Labels implied by Chan Bona (2025) |

any archived copy must carry an explicit synthetic label and pseudonymize the names of real persons. The precaution matters: writing about Philippine voter registration, Sonnet and Fable attributed invented quotations to a real election official, the readiness to put words in real mouths that Catapang et al. (2025) identified as a hazard for migrants who rely on such systems.

# 4 Results

## 4.1 Knowing

Asked to describe the modals of obligation, the models told a coherent story, and it was the literature's story (Table 2, last row). The *must* claim matches GloWbE in direction. The *need to* claim is more revealing: 12 of the 17 "less" claims concern Indian and Kenyan English, which trail the Southeast Asian varieties in Collins's ranking but have the highest *need to* rates in GloWbE; for Philippine English the claim holds, and most Singapore descriptions said *more*. The models learned Collins's conclusions but not his tables. Only 4 of 32 descriptions ventured numbers, unreliably: Sonnet placed Philippine *must* and *should* "in the low hundreds per million words in written ICE data", where Collins et al. (2014) report 777 and 1,155. Haiku, Sonnet and Opus flagged uncertainty in every response; Fable did so explicitly in none of its six consensus-coded responses.

All 31 consensus-coded descriptions asserted that the variety extends the progressive to stative verbs, which fits van Rooy (2014) and Collins (2015) but not Rautionaho (2022). Yet 27 of the 32 said the feature is rare in edited or formal writing; Opus called it "more salient than frequent" in Indian English. The models' knowledge is register-indexed.

Framing matters as much as assertion. A single "delegitimizing" code (22 of 29 consensus-coded descriptions) mixed accurate register placement with devaluation. Separated in a second, double-coded pass ($\kappa$ = .949), the pattern is sharper: of the 31 descriptions on which the coders agreed, 17 tied the variety's grammar to lower proficiency, less education, learner usage or interference (Fable 7, Opus 6, Sonnet 3, Haiku 1), as when Opus placed Singapore stative progressives in "colloquial Singlish and less educated or informal speech"; 2 others reported stigma without endorsing it; and

only 4 defended the features as systematic or locally accepted. The variety is made audible and, in the same breath, demoted: the medium audibility of Catapang (2026), now surviving a prompt that asked only for description.

Collins was named in 22 descriptions, in 44 sentences: by Opus in 8 of 8, by Sonnet and Fable in 7 of 8, and by Haiku in none, although Haiku cited Platt, Weber, Sharma, Mair and Hundt (the next most cited scholar, in 9 descriptions). Since every session named the Collins volume, how often his name appears says nothing about spontaneity; what was attributed to him can still be checked. Of the 32 sentences on which the coders agreed, 17 were supported, 12 partly supported and 3 contradicted (3 to 6 across coders): mostly right, a measure of how far the findings have traveled. But the errors lean one way. Six of the eight agreed misattributions about the progressive move Collins toward the stereotype that Outer Circle varieties overuse it, and most of the rest toward Outer Circle conservatism, as when Fable grouped Philippine and Singapore English with a "lagging" Indian English, "a pattern Collins and others have linked to 'colonial lag' ", although Collins found the Southeast Asian varieties leading (Collins, 2009) and explains such differences by evolutionary status, epicentrality and areal proximity (Collins, 2023). Two Opus sentences invert him:

(1) Corpus studies, including Collins's comparative work on the progressive across ICE corpora, suggest that Kenyan English (with Philippine and some other outer-circle varieties) has an overall progressive frequency at or above British and American levels. (Opus, Task A, Kenyan English)

Collins (2008) ranked Kenyan English, with Indian English, as the least innovative of the nine varieties in progressive use, behind British and American English. Call this *citational ventriloquism*: the model speaks in the scholar's voice, borrowing the authority of a name while severing the claim from its provenance. A human mis-citation can be traced to a page; a model's cannot, and a student asking about Kenyan English receives Collins's authority without his evidence. It is a peculiar tribute to a body of work that the machine has absorbed its author's name more securely than his numbers.

### 4.2 Doing

The first finding is what could be produced at all. Haiku declined 10 of 12 batches, explaining that realistic local news on safety and election topics could be mistaken for real reporting; told the articles were fiction for a research corpus, it completed 7 of 10 reruns but refused both Indian reruns and one British rerun, for reasons it gave as topical. The numbers cannot show a variety effect, and the other models refused none of their 36 batches. The main analysis therefore pools Sonnet, Opus and Fable (360 articles, 174,784 words). That safety behavior decides which synthetic Englishes exist is itself a methodological finding: what counts as LLM-generated "data" is shaped before any analyst touches it.

As Figure 2 shows, the Outer Circle texts carry about twice as much *must* as the Inner Circle texts, the one contrast that survives correction and clustering, and it widens when the voter-registration topic, which supplied 10 of the 15 Inner Circle tokens, is set aside. Its direction agrees with GloWbE, but its size cannot be told apart from GloWbE's ratio of 1.21. And it is a British contrast. Against American English the ratio is 1.35, close to GloWbE's 1.20; British news has three tokens of *must* in 28,911 words, all by Opus, and Haiku's British news has none in 5,074. Collins's data give no reason to expect a British–American difference (487 and 488 per million words), yet the generated American news has four times the British rate. The decline of *must* documented diachronically in the Inner Circle (Leech et al., 2009) returns as a synchronic stereotype, and it has been assigned to Britain. The models dramatize the contrast, and too eagerly.

No single narrative explains the rest. *Need to* is lower in the Outer Circle texts, as the descriptions predicted, but the contrast fails correction and resampling. *Have to* is higher, robustly to clustering though not to correction, against both the descriptions and Collins's finding that the Inner Circle leads the rise of the quasi-modals; and *should* does not differ, though 21 descriptions predicted it would. Within the Outer Circle, the *have to*:*must* ordering follows the ranking Collins (2009) made familiar, with the Southeast Asian varieties first, rather than the GloWbE ordering, but not significantly (Southeast Asian against Indian and Kenyan texts: *have to* LL = 2.52). The descriptions state Collins's ranking; the texts only echo it.

In aggregate density the instrument performs respectably: the Philippine rates are compatible with

Table 2: Obligation and necessity expressions in model-written news (Sonnet, Opus and Fable pooled), per 100,000 words (raw counts), compared with GloWbE and with the Task A descriptions

| | **Words** | ***must*** | ***have to*** | ***need to*** | ***should*** | ***have to*:*must*, generated (GloWbE)** |
|---|---|---|---|---|---|---|
| Philippine | 28,951 | 41.4 (12) | 51.8 (15) | 17.3 (5) | 62.2 (18) | 1.25 (1.49) |
| Singapore | 29,506 | 57.6 (17) | 84.7 (25) | 33.9 (10) | 57.6 (17) | 1.47 (1.95) |
| Indian | 29,021 | 55.1 (16) | 51.7 (15) | 27.6 (8) | 51.7 (15) | 0.94 (1.75) |
| Kenyan | 29,355 | 68.1 (20) | 40.9 (12) | 27.3 (8) | 74.9 (22) | 0.60 (1.33) |
| American | 29,040 | 41.3 (12) | 27.5 (8) | 37.9 (11) | 55.1 (16) | 0.67 (2.02) |
| British | 28,911 | 10.4 (3) | 38.0 (11) | 58.8 (17) | 76.1 (22) | 3.67 (2.01) |
| Outer Circle (OC) | 116,833 | 55.6 (65) | 57.3 (67) | 26.5 (31) | 61.6 (72) | |
| Inner Circle (IC) | 57,951 | 25.9 (15) | 32.8 (19) | 48.3 (28) | 65.6 (38) | |
| LL, OC vs IC (Holm-adjusted *p*) | | 8.27 (.016) | 5.10 (.069) | 5.16 (.069) | 0.10 (.758) | |
| OC:US rate ratio (exact 95% CI) | | 1.35 (0.72–2.74) | 2.08 (1.00–5.02) | 0.70 (0.34–1.54) | 1.12 (0.64–2.06) | |
| GloWbE OC:IC ratio (Collins, 2023) | | 1.21 | 0.97 | 1.05 | | |
| Task A: descriptions of the OC variety | | more: 24 of 31 | *must* preferred: 22 of 31 | less: 17 of 30 | more: 21 of 30 | |

*Note.* Ratios compare mean rates; GloWbE values indicate direction and size, not level. Without four nominal tokens (*the need to*, *no need to*), Outer Circle *need to* is 23.1.

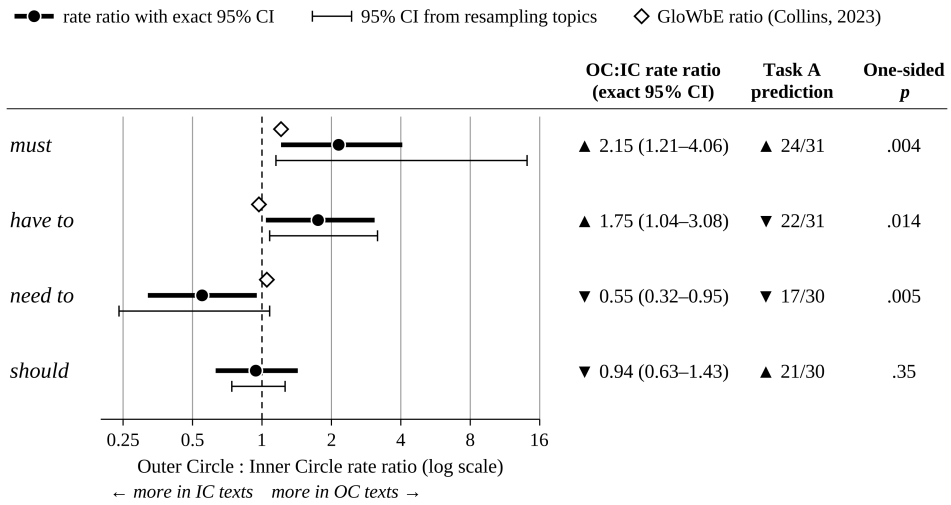



Figure 2: Outer Circle (OC) to Inner Circle (IC) rate ratios of four obligation and necessity expressions in model-written news, on a log scale. Sonnet, Opus and Fable are pooled (360 articles). Diamonds, drawn just above each bar, give the GloWbE ratio shown in Table 2. Glyphs give the direction of the observed ratio and of the prediction in the Task A descriptions (for *have to*, descriptions stating that the OC variety prefers *must*). One-sided *p* values come from a permutation test within model and topic, in the observed direction and unadjusted; the resampling intervals treat topics as the sampling unit, and the width of the interval for *must* reflects the voter-registration topic.

the PEPC for all four forms, though at this sample size little could be detected (*need to* 17.3, 95% CI 5.6–40.3, against 30.3). What the models get right is the register; what they distort is the difference between varieties, which is precisely what World Englishes research sets out to describe. Unremarked features fare unpredictably. *Shall* and *might*, at 12.6 and 15.2 per 100,000 words in the 2022 Philippine press (Dita & Chan Bona, Chapter 16, Collins volume), predict about eight tokens in the Philippine sample; none occurred (Poisson $p < .001$), although Collins et al. (2014) found *shall* more frequent in

Table 3: Agreement of model annotations with the labels implied by Chan Bona (2025)

| Model | Run 1 accuracy (95% CI) | Run 2 accuracy (95% CI) | $\kappa$ with gold (run 1 / run 2) | $\kappa$ between runs | Accuracy excluding the 16 items where gold and instructions conflict (run 1 / run 2) |
|---|---|---|---|---|---|
| Haiku | .667 (.584–.741) | .630 (.546–.706) | .510 / .457 | .915 | .748 / .706 |
| Sonnet | .667 (.584–.741) | .689 (.606–.761) | .524 / .551 | .948 | .756 / .782 |
| Opus | .741 (.661–.807) | .741 (.661–.807) | .622 / .622 | .979 | .824 / .824 |
| Fable | .696 (.614–.768) | .726 (.645–.794) | .552 / .598 | .925 | .790 / .824 |

Philippine than in British or American writing, and the Kenyan news used it three times. Philippine *may*, conversely, is inflated (103.6 English tokens per 100,000 words, against 17.2 in the American news), largely through permissions in advisories (*Voters may check their status*), which suggests a rival reading of the *must* contrast: Outer Circle news rendered as government advisory.

Register, not variety, also governs the progressive. In 224,186 words by all four models there is not one stative progressive; the broader search finds only uses standard in Inner Circle news (*are seeing more cases*, *is costing*) in all six countries, and one borderline *is requiring more complex care* (Haiku, Singapore). Given the descriptions' register caveats, this is consistency, not a gap between knowing and doing; but on this feature generated Indian news is indistinguishable from British and cannot serve as data. Localization is abundant, though partly forced by setting: *lakh(s)* 67 and *crore(s)* 45 times in the Indian texts, *matatu(s)* 24 in the Kenyan, *barangay(s)* 26 and *jeepney(s)* 14 in the Philippine, while *county* is nearly as frequent in American (94) as in Kenyan news (109). Singapore receives *hawker* 19 times and no Singlish. Only the Philippine texts go beyond nouns: 45 of their 60 articles quote speakers in Tagalog or Taglish, and no other variety is quoted in a local language. The matrix grammar stays standard:

(2) Local government units across Bicol and Eastern Visayas began moving thousands of residents out of coastal and landslide-prone barangays on Tuesday as Typhoon Nonoy continued to gain strength over the Philippine Sea. (Opus, Task B, Philippine English)

### 4.3 Annotating

All four models were highly consistent with themselves and only moderately consistent with the analyst (Table 3).

Opus was numerically the most accurate, but after correction for 12 pairwise comparisons only its second-run advantage over Haiku is significant (adjusted $p = .009$). A majority vote over all eight runs scored .689, below Opus alone, because it averages in weaker models and the runs share their errors. Between-run $\kappa$ of .915 to .979 alongside gold-standard $\kappa$ of .457 to .622 describes an instrument that is reliably unreliable: repetition will not reveal its mistakes, because it repeats them.

The errors are not random. Of the 28 items that all eight runs got wrong, 14 are items on which the gold contradicts the instructions the models were given, which offered *He wants to run* and *The team needs to win to qualify* as dynamic and reserved "other" for past-time uses: six instances of *want(s) to*, two of *need to* and six temporal uses of *would*, *was to* and *was about to*, which the thesis reads otherwise (Table 3, last column). The remainder include dynamic *may* and *might*, preterite *should* and deontic *will*, where a trained reading departs from a textbook default, and a few debatable gold assignments (*This is to give complete justice to the victim*, read as epistemic). Agreement in annotation depends on whose definitions the annotator is given, the researcher's own included. The gold comes from a single analyst, and examples chosen to illustrate categories are probably easier than running data; the contrast with the more than 90% that Morin and Marttinen Larsson (2025) achieved on a binary formal variable is a reminder that what is annotated matters at least as much as which model annotates it.

## 5 Discussion: instrument, object, agent

### 5.1 Instrument

As an instrument, the model is most useful where validation is cheapest and least useful where description is hardest. Describing, it returned the secondary literature's consensus under an authority it could not cite reliably; producing, it exaggerated some contrasts and dropped unremarked features; annotating, it repeated its errors. These are not corrigible glitches but structural features of a system that predicts coherence, not truth, and they suggest why, in a preprint replicating 37 annotation tasks, LLM labels led to incorrect conclusions for roughly a third of the hypotheses (Baumann et al., 2025). The reframing this requires is simple to state and demanding to practice: every generated artifact is a proposal that solicits verification, not an observation that demands explanation. The Task A descriptions served as checkable hypotheses, two of which held in production and two of which failed. Their divergences from the corpora are data of a second kind, a record of what the wider textual public knows, or believes it knows, about a variety; that reputational knowledge is itself sociolinguistic data about visibility and stereotype. Model labels, likewise, acquire evidential value only once a human-coded subsample locates their error, and the definitions the model received belong in the report.

### 5.2 Object

As an object, model-written English deserves description in its own right. It is a variety in the representational sense of Grieve et al. (2025), a model of the text it was trained on, but not in the sociolinguistic sense, having no speakers and no community. More exactly, it is a single edited register onto which varieties are drawn as lexical overlays, a post-variety in Seargeant and Tagg's (2011) sense, and perhaps the purest one. Its rule is frequency for registers and salience for varieties, which is why the Philippine densities match the press while contrasts between varieties are distorted. In Labov's (1972) terms, the models reproduce stereotypes, such as the Outer Circle's attachment to *must*, and lose indicators, such as Philippine *shall*, much as human speakers do when asked to perform a variety (Schilling-Estes, 1998). One mechanism is plausible: a prompt that names a country retrieves what is associated with the name, and names attach mainly to text *about* varieties, since text *in* them is seldom labeled, so that the model learns registers from the latter and varieties from the former. Preference tuning (Ryan et al., 2024) and stereotypes of advisory journalism compete as explanations, and proprietary models do not allow pretraining and post-training effects to be separated; open-weight models, and prompts conditioned on attested local text rather than a name, are the obvious tests. Read against the corpora, such output becomes a probe that, when pressed against the field's evidentiary surface, reveals where the surface is thin.

### 5.3 Agent: algorithmic exonormativity

As an agent, the model matters because its English does not stay in the laboratory. Chatbots now take on tasks that migrants once handled through human interaction (Catapang et al., 2025), AI writing suggestions shift Indian writers toward Western styles (Agarwal et al., 2025), and models recognize Asian Englishes without authorizing them (Catapang, 2026). That study suggested that explicit prompting could undo such defaults; the present results qualify the suggestion. Situated prompting recovers the lexicon, and for the Philippines quoted Tagalog, but not the grammar or the frequencies, and descriptions that recognize local grammar still tie it to low proficiency.

This paper proposes the term *algorithmic exonormativity* for the condition in which the defaults of globally deployed models acquire normative force in a community's writing: returned as corrections, offered as models and issued as metalinguistic judgments. The data document its first half, an external norm on offer; whether communities orient to it, through accepted rewrites, attitudes or editorial and teaching guidelines, is a question the paper can frame but not answer. Its timing is the irony. Collins et al. (2014) and Borlongan (2016) place Philippine English at or near endonormative stabilization, which Schneider (Chapter 1, Collins volume) treats as a composite of political independence, attitude (in his estimation the strongest factor), homogeneity, structural distinctness, codification and literary creativity, and in which he sees no clear signs of reversal. The models bear directly on three of these components (Figure 3): attitude, by placing local grammar on a scale of proficiency; codification, by reproducing the asymmetry Schneider describes, local words admitted and local

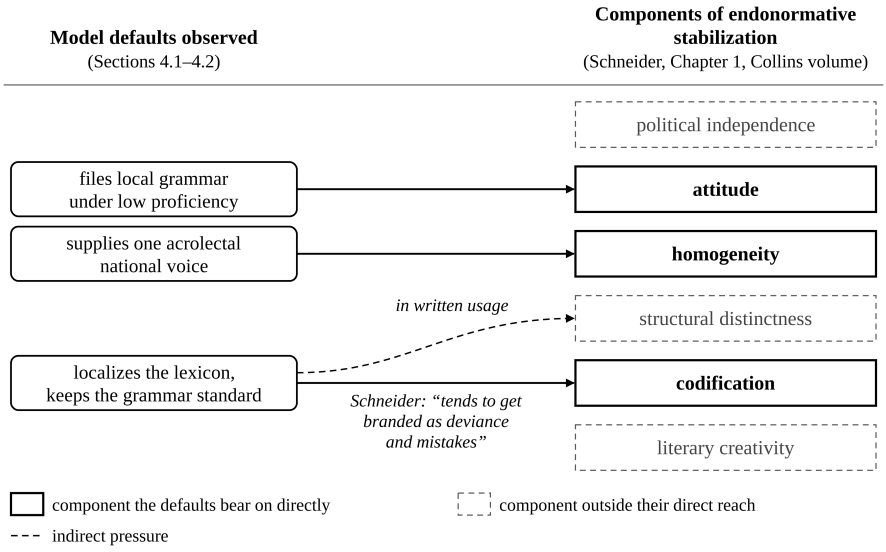



Figure 3: Model defaults observed in Tasks A and B mapped onto Schneider's components of endonormative stabilization (Chapter 1, Collins volume). Schneider hypothesizes that attitude is the strongest of them; the quotation is his description of how localized grammar is received.

grammar branded as deviance; and homogeneity, by supplying a single acrolectal national voice. Structural distinctness is touched only indirectly, through the written usage the models would standardize; political independence lies beyond what a writing assistant supplies or withholds; and literary creativity, which such assistants may well touch, was not examined here. Because local words are admitted and local grammar is not, the pressure could arrest stabilization at its lexical stage; because it leaves the other components standing, it would not reverse the Dynamic Model (Schneider, 2007). In Collins's terms it is also anti-colloquializing: the features it files as informal, extended progressives first among them, are those that colloquialization carries into writing (Collins, 2008; Collins & Yao, 2013).

The new center differs from the colonial one. Placeless in self-presentation, it still has an address near American written English (Ryan et al., 2024), and it stereotypes even British English, pointing to a norm that nobody speaks. Unlike spell-checkers and house styles, it generates, and it explains with citations, so that part of the center is World Englishes scholarship itself, returned to its subjects as prescription. In the image of Catapang (2026), the hourglass does not flatten the circles; it reorders their authority. Buschfeld (Chapter 2, Collins volume) asks whether computer-mediated communication is an extra- or an intra-territorial force; this one is both, built outside the territory and used inside local acts of composition. Her multinormative phase, in which local, regional and global standards coexist (Meer & Deuber, 2020), is the strongest counter-reading, but multinormativity describes which norms are available, algorithmic exonormativity which norm is present by default. The Extra- and Intra-territorial Forces model allows development to reverse (Buschfeld & Kautzsch, 2017), which gives a re-emerging external norm a theoretical place; and unlike transnational attraction (Schneider, 2014), in which speakers appropriate English as a resource, here a norm is supplied to them.

Nothing in the technology fixes the direction. A model tuned on local corpora could serve codification, and writers can ignore, resist or exploit what is offered, the "standardisation paradox" of AI that both homogenizes and pluralizes (Ugwuanyi et al., 2025). The costs, meanwhile, are not evenly borne. In classrooms, job applications and immigration paperwork, where dialect is already an axis of discrimination (Hofmann et al., 2024), an assistant that files local grammar under low proficiency can make vulnerable speakers more vulnerable still.

### 5.4 The post-2022 corpus problem

The agent relation also reaches back into the evidence. At least 13.5% of biomedical abstracts published in 2024 were processed with LLMs (Kobak et al., 2025), and a widely used word-frequency resource was frozen at 2021 data partly because the web "is full of slop generated by large language models" (Speer, 2024). Models trained recursively on generated text lose the tails of the original distribution (Shumailov et al., 2024), and minority-variety features may be such tails. No study known to this author has quantified machine-generated text in the web corpora on which World Englishes research relies. Table 2 predicts its signature: wider salient contrasts (more Outer Circle *must* and less Outer Circle *need to*, relative to the Inner Circle), an exaggerated British retreat from *must*, and the erasure of *shall* where it survives. Change driven by Americanization or colloquialization should instead bring the circles closer, so the two are distinguishable in principle. Not all of this is contamination: where writers accept model rewrites, the text is human and its shifts are change, which makes machine mediation a candidate driver beside Americanization and colloquialization, testable with matched corpora. Corpora with documented provenance, the signature of most of Collins's work, gain value, and pre-2022 web corpora such as the GloWbE on which his recent comparisons rest become baselines that no later web sample can simply replicate. So do resources such as the social media data of Gonzales (Chapter 5, Collins volume) and the 2022 press corpus of Dita and Chan Bona (Chapter 16, Collins volume); their successors will need to record not only who wrote a text, but whether anyone did.

## 6 Conclusion

Asked to describe, to write and to annotate English worldwide, the models tested here returned, each time, a version of what the field has said rather than what its corpora show, and even that only in part. Deployed at scale, that version is positioned to become a norm: an exonormative center whose influence on usage, and on the corpora that record usage, the field has barely begun to measure.

The limits deserve equal clarity: one vendor's models on one day, in a context that named the honorand; one wording per task, although trivial changes of format can move model behavior (Sclar et al., 2024), and a Task A wording that presupposed the stative feature; one register and one country's news agenda; small counts; codes from instances of a model under test; and one analyst's gold labels on illustrative examples. Each limit points to a next step: other vendors and open-weight models; prompts that name the variety, condition on attested text or move to informal registers; balanced topics; human multi-coder annotation of running data; and, above all, judgments by speakers of the varieties concerned, the only arbiters of whether a synthetic text sounds like home or like a caricature of it.

Table 4 condenses the lessons into a reporting checklist, drawing on Törnberg (2024), Chen et al. (2024) and Palmer et al. (2024). Its purpose is modest: to make the model a variable to be analyzed rather than an invisible background instrument. Its larger aim is a different music from the model's confident cadence, one that makes uncertainty audible, foregrounds the work of checking and keeps local voices in front.

The direction of the argument does not depend on any one tier of the family tested. Every tier wrote local vocabulary into unmarked grammar; every tier that cited Collins misstated him at least in part, though the smallest errors, such as a co-author borrowed from a companion paper, were Sonnet's; and every tier repeated its annotation errors. The largest model was numerically the most accurate annotator, but no tier was, in any way these tasks could detect, closer to the varieties themselves. Collins taught the field to ask of every claim what the corpus shows. Put to these systems, that question has an uncomfortable answer: they show us ourselves, our literature, our narratives, our names. Until evaluation and training change, the appropriate posture is disciplined use: let the model propose, let the corpora and communities decide, and let the record show not only what English looks like in the world, but also how our instruments try to reshape it as they look.

Table 4: Reporting checklist for LLM use in World Englishes research

| Item | What to report |
|---|---|
| Model and date | Exact model identifier, access route and date of every run; models change silently under stable names |
| Settings and context | Sampling settings, tools, system prompt, session or project context and anything else the model received besides the prompt |
| Prompts | Verbatim prompts, including framings added after refusals, and whether the variety or its features were named |
| Repetition | Number of independent runs and run-to-run agreement, reported separately from agreement with gold labels |
| Refusals | Refusals and non-responses by model, condition and topic |
| Benchmarks | Named corpora, register match, topic balance, normalization basis and uncertainty (confidence intervals, clustering, correction for multiple comparisons) |
| Coders and gold standard | Who or what coded, whether coders share a model family with the systems tested, whose labels serve as gold, which framework, and how definitions and examples were worded |
| Provenance | Synthetic text labeled, never pooled with attested data, real names pseudonymized, archived with prompts and scripts |

## References


Agarwal, D., Naaman, M., & Vashistha, A. (2025). AI suggestions homogenize writing toward Western styles and diminish cultural nuances. In *Proceedings of the 2025 CHI Conference on Human Factors in Computing Systems* (pp. 1–21). Association for Computing Machinery. https://doi.org/10.1145/3706598.3713564

Baumann, J., Röttger, P., Urman, A., Wendsjö, A., Plaza-del-Arco, F. M., Gruber, J. B., & Hovy, D. (2025). *Large language model hacking: Quantifying the hidden risks of using LLMs for text annotation* (arXiv:2509.08825) [Preprint]. arXiv. https://arxiv.org/abs/2509.08825

Borlongan, A. M. (2016). Relocating Philippine English in Schneider's dynamic model. *Asian Englishes, 18*(3), 232–241. https://doi.org/10.1080/13488678.2016.1223067

Buschfeld, S., & Kautzsch, A. (2017). Towards an integrated approach to postcolonial and non-postcolonial Englishes. *World Englishes, 36*(1), 104–126. https://doi.org/10.1111/weng.12203

Catapang, J. K. (2026). Asymmetrical pluralism and the normative power of artificial intelligence in Asian Englishes. *Asian Englishes*. Advance online publication. https://doi.org/10.1080/13488678.2026.2685377

Catapang, J. K., Borlongan, A. M., & Go, M. A. C. (2025). Language, migration, and ChatGPT. *Journal of Modern Languages, 35*(2), 167–189. https://doi.org/10.22452/jml.vol35no2.9

Chan Bona, K. F. U. (2025). *Modal auxiliary verbs in Philippine English press: A corpus-based analysis* [Master's thesis, De La Salle University]. Animo Repository. https://animorepository.dlsu.edu.ph/etdm_deal/23

Chen, L., Zaharia, M., & Zou, J. (2024). How is ChatGPT's behavior changing over time? *Harvard Data Science Review, 6*(2). https://doi.org/10.1162/99608f92.5317da47

Collins, P. (2008). The progressive aspect in World Englishes: A corpus-based study. *Australian Journal of Linguistics, 28*(2), 225–249. https://doi.org/10.1080/07268600802308782

Collins, P. (2009). Modals and quasi-modals in world Englishes. *World Englishes, 28*(3), 281–292. https://doi.org/10.1111/j.1467-971X.2009.01593.x

Collins, P. (2015). Recent diachronic change in the progressive in Philippine English. In P. Collins (Ed.), *Grammatical change in English world-wide* (pp. 271–296). John Benjamins. https://doi.org/10.1075/scl.67.12col

Collins, P. (2023). Modals and quasi-modals in English world-wide. *Journal of English Linguistics, 51*(3), 265–293. https://doi.org/10.1177/00754242231173720

Collins, P., Borlongan, A. M., & Yao, X. (2014). Modality in Philippine English: A diachronic study. *Journal of English Linguistics, 42*(1), 68–88. https://doi.org/10.1177/0075424213511462

Collins, P., & Yao, X. (2013). Colloquial features in World Englishes. *International Journal of Corpus Linguistics, 18*(4), 479–505. https://doi.org/10.1075/ijcl.18.4.02col

Davies, M., & Fuchs, R. (2015). Expanding horizons in the study of World Englishes with the 1.9 billion word Global Web-based English Corpus (GloWbE). *English World-Wide, 36*(1), 1–28. https://doi.org/10.1075/eww.36.1.01dav

Fleisig, E., Smith, G., Bossi, M., Rustagi, I., Yin, X., & Klein, D. (2024). Linguistic bias in ChatGPT: Language models reinforce dialect discrimination. In Y. Al-Onaizan, M. Bansal, & Y.-N. Chen (Eds.), *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing* (pp. 13541–13564). Association for Computational Linguistics. https://doi.org/10.18653/v1/2024.emnlp-main.750

Gilardi, F., Alizadeh, M., & Kubli, M. (2023). ChatGPT outperforms crowd workers for text-annotation tasks. *Proceedings of the National Academy of Sciences, 120*(30), Article e2305016120. https://doi.org/10.1073/pnas.2305016120

Grieve, J., Bartl, S., Fuoli, M., Grafmiller, J., Huang, W., Jawerbaum, A., Murakami, A., Perlman, M., Roemling, D., & Winter, B. (2025). The sociolinguistic foundations of language modeling. *Frontiers in Artificial Intelligence, 7*, Article 1472411. https://doi.org/10.3389/frai.2024.1472411

Hofmann, V., Kalluri, P. R., Jurafsky, D., & King, S. (2024). AI generates covertly racist decisions about people based on their dialect. *Nature, 633*(8028), 147–154. https://doi.org/10.1038/s41586-024-07856-5

Huddleston, R., & Pullum, G. K. (2002). *The Cambridge grammar of the English language*. Cambridge University Press. https://doi.org/10.1017/9781316423530

Kachru, B. B. (1985). Standards, codification and sociolinguistic realism: The English language in the Outer Circle. In R. Quirk & H. G. Widdowson (Eds.), *English in the world: Teaching and learning the language and literatures* (pp. 11–30). Cambridge University Press.

Kobak, D., González-Márquez, R., Horvát, E.-Á., & Lause, J. (2025). Delving into LLM-assisted writing in biomedical publications through excess vocabulary. *Science Advances, 11*(27), Article eadt3813. https://doi.org/10.1126/sciadv.adt3813

Labov, W. (1972). *Sociolinguistic patterns*. University of Pennsylvania Press.

Leech, G., Hundt, M., Mair, C., & Smith, N. (2009). *Change in contemporary English: A grammatical study*. Cambridge University Press. https://doi.org/10.1017/CBO9780511642210

Meer, P., & Deuber, D. (2020). Standard English in Trinidad: Multinormativity, translocality, and implications for the Dynamic Model and the EIF Model. In S. Buschfeld & A. Kautzsch (Eds.), *Modelling World Englishes: A joint approach to postcolonial and non-postcolonial varieties* (pp. 274–297). Edinburgh University Press. https://doi.org/10.1515/9781474445887-016

Morin, C., & Marttinen Larsson, M. (2025). Large corpora and large language models: A replicable method for automating grammatical annotation. *Linguistics Vanguard, 11*(1), 501–510. https://doi.org/10.1515/lingvan-2024-0228

Palmer, A., Smith, N. A., & Spirling, A. (2024). Using proprietary language models in academic research requires explicit justification. *Nature Computational Science, 4*(1), 2–3. https://doi.org/10.1038/s43588-023-00585-1

Rautionaho, P. (2022). Revisiting the myth of stative progressives in world Englishes. *World Englishes, 41*(2), 183–206. https://doi.org/10.1111/weng.12520

Reinhart, A., Markey, B., Laudenbach, M., Pantusen, K., Yurko, R., Weinberg, G., & Brown, D. W. (2025). Do LLMs write like humans? Variation in grammatical and rhetorical styles. *Proceedings of the National Academy of Sciences, 122*(8), Article e2422455122. https://doi.org/10.1073/pnas.2422455122

Ryan, M. J., Held, W., & Yang, D. (2024). Unintended impacts of LLM alignment on global representation. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (pp. 16121–16140). Association for Computational Linguistics. https://doi.org/10.18653/v1/2024.acl-long.853

Schilling-Estes, N. (1998). Investigating "self-conscious" speech: The performance register in Ocracoke English. *Language in Society, 27*(1), 53–83. https://doi.org/10.1017/S0047404500019722

Schneider, E. W. (2007). *Postcolonial English: Varieties around the world*. Cambridge University Press. https://doi.org/10.1017/CBO9780511618901

Schneider, E. W. (2014). New reflections on the evolutionary dynamics of world Englishes. *World Englishes, 33*(1), 9–32. https://doi.org/10.1111/weng.12069

Sclar, M., Choi, Y., Tsvetkov, Y., & Suhr, A. (2024). Quantifying language models' sensitivity to spurious features in prompt design or: How I learned to start worrying about prompt formatting. In *The Twelfth International Conference on Learning Representations (ICLR 2024)*. https://iclr.cc/virtual/2024/poster/18650

Seargeant, P., & Tagg, C. (2011). English on the internet and a "post-varieties" approach to language. *World Englishes, 30*(4), 496–514. https://doi.org/10.1111/j.1467-971X.2011.01730.x

Shumailov, I., Shumaylov, Z., Zhao, Y., Papernot, N., Anderson, R., & Gal, Y. (2024). AI models collapse when trained on recursively generated data. *Nature, 631*(8022), 755–759. https://doi.org/10.1038/s41586-024-07566-y

Speer, R. (2024, September). *Why wordfreq will not be updated*. GitHub. https://github.com/rspeer/wordfreq/blob/master/SUNSET.md

Törnberg, P. (2024). Best practices for text annotation with large language models. *Sociologica, 18*(2), 67–85. https://doi.org/10.6092/issn.1971-8853/19461

Uchida, S. (2024). Using early LLMs for corpus linguistics: Examining ChatGPT's potential and limitations. *Applied Corpus Linguistics, 4*(1), Article 100089. https://doi.org/10.1016/j.acorp.2024.100089

Ugwuanyi, K., Mair, C., Dovchin, S., Erdocia, I., Kuteeva, M., & Airemionkhale, E. (2025). Generative AI and linguistic diversity in academic writing and publishing: Perspectives from World Englishes. *Journal of English for Research Publication Purposes, 6*(2), 265–289. https://doi.org/10.1075/jerpp.00035.ugw

van Rooy, B. (2014). Progressive aspect and stative verbs in Outer Circle varieties. *World Englishes, 33*(2), 157–172. https://doi.org/10.1111/weng.12079